



Google Cloud

Deriving Insights from Unstructured Data using Machine Learning

Carolyn Ujcic

Cloud Machine Learning Engineer

Hi I'm Carolyn Ujcic, a Cloud Machine Learning Engineer. So far in this course you've been applying machine learning to your structured datasets with custom models. It's time to widen our aperture and look at unstructured datasets (like images) and also other approaches to doing ML that don't involve custom model building.

Agenda

ML drives business value

How does ML on unstructured data work?

Choosing the right ML approach

- Pre-built AI building blocks
- Using Pre-built AI to create a chatbot
- Customizing Pre-built models with AutoML
- Building a custom model

Demo: Text classification done three ways

Here's our agenda. We'll first see how ML on unstructured datasets is driving value for businesses.

Agenda

ML drives business value

How does ML on unstructured data work?

Choosing the right ML approach

- Pre-built AI building blocks
- Using Pre-built AI to create a chatbot
- Customizing Pre-built models with AutoML
- Building a custom model

Demo: Text classification done three ways

Then we'll look at the intuition and models behind unstructured data

Agenda

ML drives business value

How does ML on unstructured data work?

Choosing the right ML approach

- Pre-built AI building blocks
- Using Pre-built AI to create a chatbot
- Customizing Pre-built models with AutoML
- Building a custom model

Demo: Text classification done three ways

After that is a critical topic on choosing the right ML approach: whether to build a model from scratch or use pre-existing building blocks.

Agenda

ML drives business value

How does ML on unstructured data work?

Choosing the right ML approach

- Pre-built AI building blocks
- Using Pre-built AI to create a chatbot
- Customizing Pre-built models with AutoML
- Building a custom model

Demo: Text classification done three ways

Lastly - we'll do a demo to classify text using three different approaches and compare the results.

Examples of real-world ML use cases



Here are just a few customer examples of companies who have used the ML tools on Google Cloud Platform on their datasets. The left-to-right ordering is again reflective of more low-level programming to more abstraction and UI-based.

Examples of real-world ML use cases

Custom image
model to price
cars



Aucnet built their own custom model to classify images of car parts and estimate price.

Examples of real-world ML use cases

Custom image
model to price
cars

Build off NLP
API to route
customer emails



Ocado used parsed results from the NL API to route customer emails to the correct responders.

Examples of real-world ML use cases

Custom image
model to price
cars

Build off NLP
API to route
customer emails

Use Vision API
as-is to find text
in memes



Giphy uses the out-of-the-box Vision API to find the text in memes using optical character recognition. It then can reject inappropriate uploads based on sentiment or keywords.

Examples of real-world ML use cases

Custom image
model to price
cars

Build off NLP
API to route
customer emails

Use Vision API
as-is to find text
in memes

Use Dialogflow
to create a new
shopping
experience



Uniqlo (uni-clo) designed a shopping chatbot using A Dialogflow UI. Dialogflow is a google-owned company which specializes in building ML-based interfaces like intelligent chatbots.

Let's look at some other use cases for ML in business

[Resource]

Google NeXT Customer Success Stories

https://youtu.be/BwWg_HVfsM?t=4m41s

Clouds or snow-capped mountains?



You might be Airbus, and use machine learning to differentiate between clouds and snow cover. If you're stumped like I am, the clouds are in the upper-right part of the right image highlighted in red.

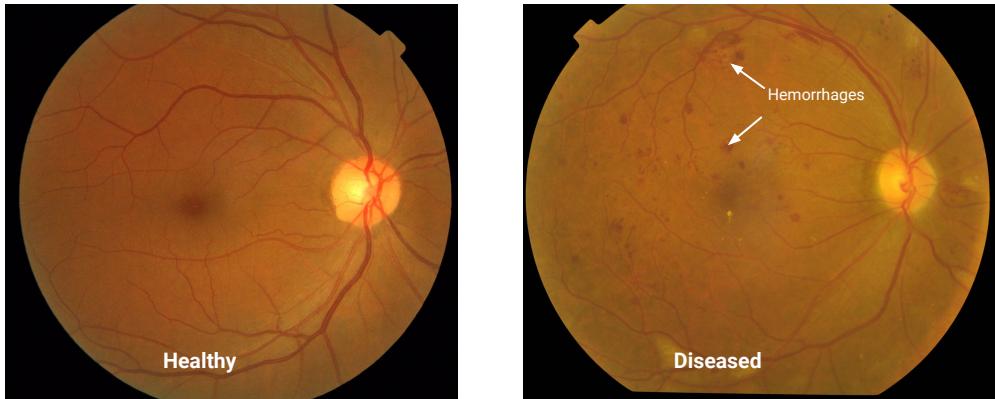
<https://cloud.google.com/blog/products/gcp/google-cloud-machine-learning-now-open-to-all-with-new-professional-services-and-education-programs>

Empty or Full?



You might be an economic forecast firm looking to track the global fleet of container ships via satellite imagery -- knowing the amount of cargo being carried might help improve your economic forecast days or months ahead of the official numbers.

Diagnosing Diabetic Retinopathy



Medical images are ripe for innovation. For example, You could diagnose [medical conditions](#) like Diabetic Retinopathy earlier when it's easier to treat and prevent blindness.

<https://ai.googleblog.com/2016/11/deep-learning-for-detection-of-diabetic.html>



Recap: Image classification
automates tasks that are easy
(and not easy) for humans



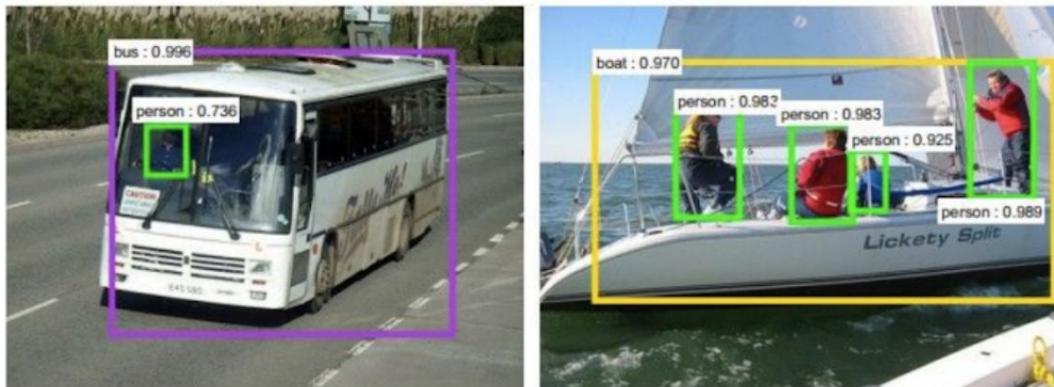
The key takeaway is that ML can automate tasks that may save or assist a human team and, recently, the latest models are even out-performing humans in some domains.

Machine learning can label images or video



Also, image classification as a field is more than just a binary classification tool. Later in this course you'll experiment with the pre-built Vision API model, which allows you to pass through a JSON request and get back a ranked list of associated labels for your image.

Machine learning can detect objects in images or video



Also, if you have more than one subject matter in a photo, it can draw bounding boxes and classify pieces of an image as well.

Describes without errors	Describes with minor errors	Somewhat related to the image	Unrelated to the image
			
A person riding a motorcycle on a dirt road.	Two dogs play in the grass.	A skateboarder does a trick on a ramp.	A dog is jumping to catch a frisbee.
			
A group of young people playing a game of frisbee.	Two hockey players are fighting over the puck.	A little girl in a pink hat is blowing bubbles.	A refrigerator filled with lots of food and drinks.
			
A herd of elephants walking across a dry grass field.	A close up of a cat laying on a couch.	A red motorcycle parked on the side of the road.	A yellow school bus parked in a parking lot.

Show and Tell: A Neural Image Caption Generator Vinyals et al 2015:
<https://arxiv.org/abs/1411.4558>

Modern image classification models can even generate captions describing what is going on in the image, like a map of dependencies like “two hockey players are fighting over a puck”. Here, it’s important to call out that even the best models can and will make mistakes in their predictions (like the road sign captioned as “a refrigerator filled with lots of food and drinks”.)

Having fun with ML: Pose Detection



g.co/movemirror

As you saw, there were a lot of impressive uses for machine learning these days, like [detecting objects in images](#), helping to [detect diseases](#), and even enabling [cars to drive themselves](#).

But AI can also be used in more playful ways too. Through a pose-estimation model, a Google AI Experiment called Move Mirror can match your real-time movements to hundreds of images of people doing similar poses from around the world. Feel free to try it out yourself and have some fun. Then, tune back into this course to learn how image classification models extract features like these from images.

<https://www.blog.google/technology/ai/move-mirror-you-move-and-80000-images-move-you/>

<https://experiments.withgoogle.com/move-mirror>

<https://medium.com/tensorflow/move-mirror-an-ai-experiment-with-pose-estimation-in-the-browser-using-tensorflow-js-2f7b769f9b23>

Agenda

ML drives business value

How does ML on unstructured data work?

Choosing the right ML approach

- Pre-built AI building blocks
- Using Pre-built AI to create a chatbot
- Customizing Pre-built models with AutoML
- Building a custom model

Demo: Text classification done three ways

Earlier in the course, you built custom models with BigQuery ML using columns of data as features and predicted the value of your labels. But what about unstructured data? How does ML learn the features?



It's best illustrated with a quick example.

What do you see here? It's a cat -- but how you know?

We know this is a cat, but how would
you teach a machine?



Your eyes have the benefit of many many years of evolution and intuition to allow you to perceive and interpret those pixels on the screen. How could we teach a machine to understand that this particular collection of pixels is a cat?

We know this is a cat, but how would
you teach a machine?

Maybe we
highlight the
eyes?



If you let yourself fall back into rules making, you might say look for cat-like eyes in the images.

What about this?

Okay...how
about eyes
OR ears



What about this image? Your brain still knows it's likely a cat, but the machine now has no basis to go off of with our old rule of "look at the eyes".

We've added more *rules* that we pre-define as "cat-like"



Okay, what if we added a bunch more hardcoded rules, like look for eyes, ears, and a nose?

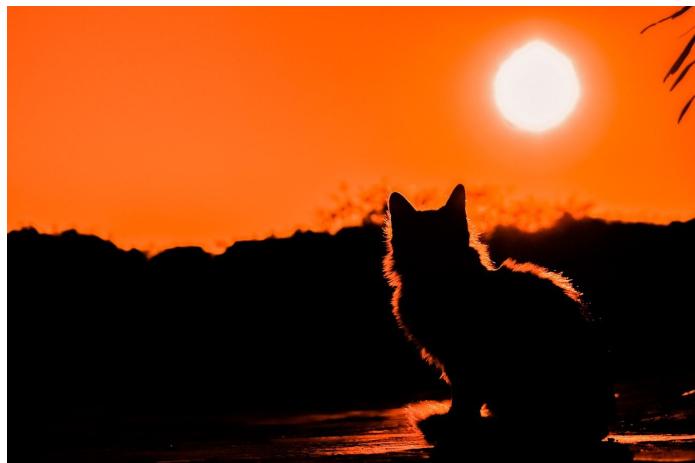
What about this?

Scores pretty
low on our
rule for cat
ears



So, is this still a cat?

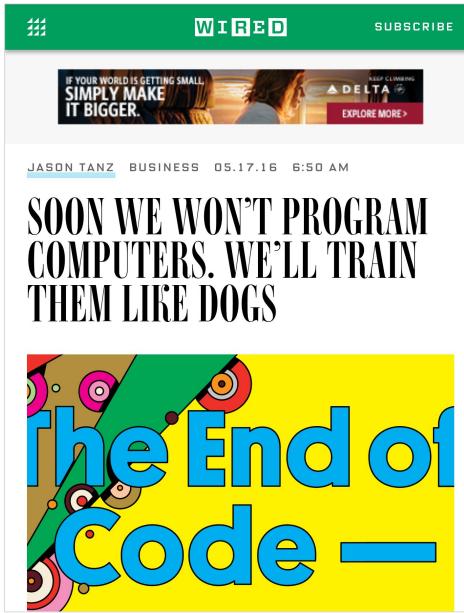
Or even this?



No eyes are
even present.
This is hard!

What about this?

Again, hardcoded rules completely fails us here and that's where deep learning comes into play.



WIRED's headline

"If you want to teach a neural network to recognize a cat, for instance, you don't tell it to look for whiskers, ears, fur, and eyes. You simply show it thousands and thousands of photos of cats, and eventually it works things out."

As this Wired article states, if you want to teach a neural network to recognize a cat you simply show it thousands of photos of cats and let it figure it out.

This is similar to how we teach children to recognize and classify new objects.

Google in 2012: Show the computer 10 million images, have it find cats



And in 2012 that's exactly what the Google Research team with Jeff Dean and Andrew Ng did. What you see here is what the deep learning neural network figured out what a cat is based on looking at over 10 million images and processing the model over 16,000 computers.

[resource]

2012 Research Paper 16,000 computers in a neural network to identify a cat face

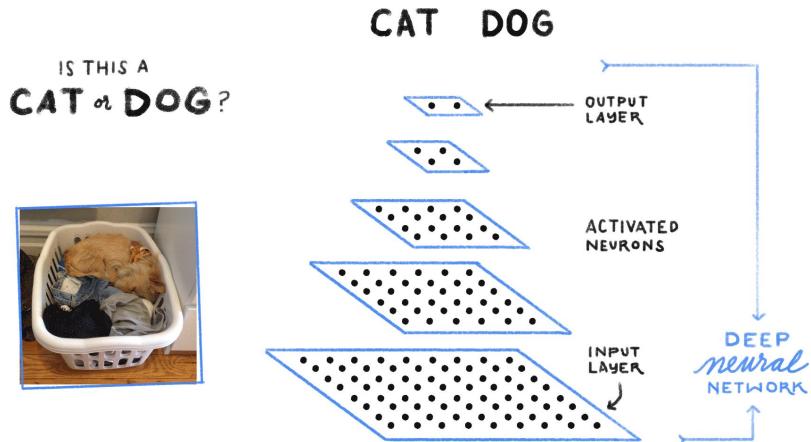
<https://static.googleusercontent.com/media/research.google.com/en//pubs/archive/38115.pdf>

Follow Google's AI research blog:

<https://ai.googleblog.com/>

<https://googleblog.blogspot.com/2012/06/using-large-scale-brain-simulations-for.html>

Modern AI Applications use Deep Learning



Now a familiar architecture for deep learning is the neural network which is a model inspired by our own human brain. Here it takes the input image and classifies it as a cat or dog.

Again, we're not telling the model to focus on looking for dog collars or cat whiskers, it builds its own recipe for determining the correct label to apply at the end.

And, as you can see from the image, modern ML models can scale and handle even tricky data points like this dog hiding in a laundry basket!

Google NN (gif source):

<http://selmandesign.com/qa-on-machine-learning/>



Machine Learning (specifically Deep Learning) is at the core of many Google Products like Google Photos, which can classify and group photos (like photos of your pets) together in an album.

Note that this is now multiple Machine Learning models in one:

- First it needs to identify whether the photo you just took is of a dog
- Second it needs to identify whether this dog is your pet based on comparing this photo to the history of photos you have with that specific dog

[resource]

Google Photos (gif source):

<https://www.blog.google/products/photos/meow-its-even-easier-find-your-furry-friends-google-photos/>

Google Photos Blog:

<https://www.blog.google/products/photos/>

Use Deep Learning when you can't explain the labeling rules

Deep Learning (remember that's a sub-discipline of Machine Learning) is incredibly useful when we as humans can't even map out our intuition about what makes a prediction correct or not. That's where we just show the computer ten thousand images of a cat and hope it figures it out. And with recent leaps in ML research at Google, computer vision has come a very long way.

Agenda

ML drives business value

How does ML on unstructured data work?

Choosing the right ML approach

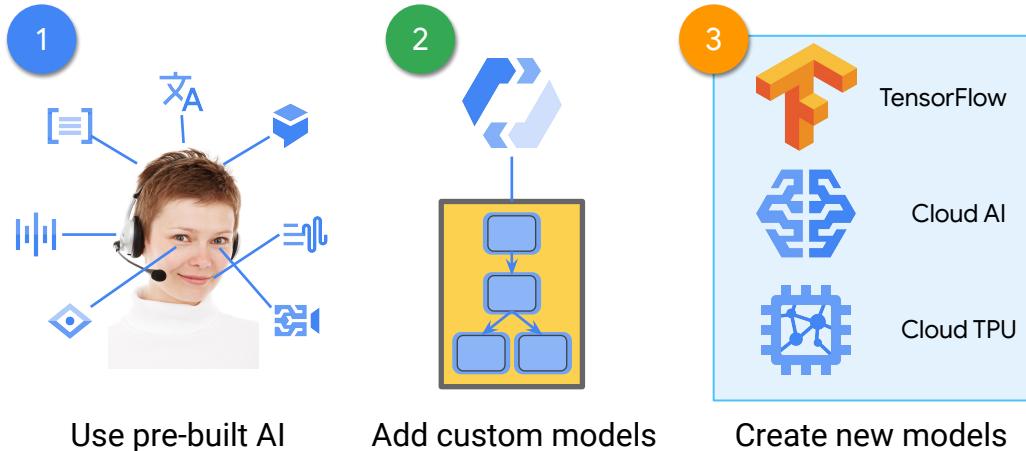
- Pre-built AI building blocks
- Using Pre-built AI to create a chatbot
- Customizing Pre-built models with AutoML
- Building a custom model

Demo: Text classification done three ways

We've covered machine learning throughout this course but that was just one approach -- custom model building. In our first challenge you may recall you inherited a Spark ML job that your data science team created to predict housing rentals. The second challenge had you create a forecasting model with BigQuery ML from scratch -- although you saved time by coding and running it in just SQL.

It's now time to zoom back out on approaches to machine learning and even look at some ways we can apply it without having to write code at all.

Artificial Intelligence application strategy



There are three approaches to AI that you should consider. You have already seen and built custom models with BQML and we have a separate set of courses on TensorFlow for even deeper model building. A good rule of thumb is to consider custom model building only when you have a lot of data like 100,000+ to millions of examples.

But what if you don't? Consider using pre-built AI which are models like the Video Intelligence and Cloud Vision APIs that you saw before. In addition, if you're looking to build a chatbot start with DialogFlow, which is a full-fledged application with ML built-in.

But what about if the building blocks don't work well for the specificity you need on your data? That's when you consider AutoML as a good candidate. It can even work with just a little bit of data like 10-100 images per label.

This lesson covers each of these approaches in detail.

Agenda

ML drives business value

How does ML on unstructured data work?

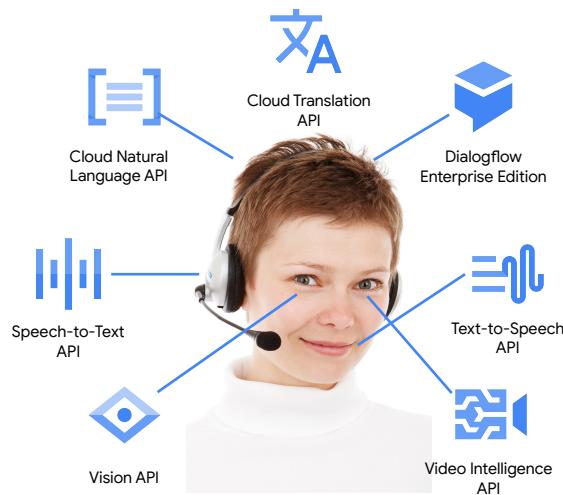
Choosing the right ML approach

- Pre-built AI building blocks
 - Using Pre-built AI to create a chatbot
 - Customizing Pre-built models with AutoML
 - Building a custom model

Demo: Text classification done three ways

First up is using pre-built AI building blocks for your use case

Do you need a custom model?



Pre-built models are offered as services. In many cases these building blocks can be used to create the application you want without the expense or complexity of creating your own models.

Cloud Speech-to-Text converts audio to text for data processing. Cloud Natural Language API recognizes parts of speech called entities and sentiment. Cloud Translation converts text in one language to another. Dialogflow Enterprise Edition is used to build chatbots to conduct conversations. Cloud Text-to-Speech converts text into high quality voice audio. Cloud Vision API is for working with and recognizing content in still images. And Cloud Video Intelligence API is for recognizing motion and action in video.

Good Machine Learning models require **lots of high-quality training data**

Good Machine Learning models require lots of high-quality training data. As we mentioned before, you should aim for 100,000+ records to train on for a custom model. If you don't have that kind of data, pre-built models are a great place to start.

Let's take a look at one of the most popular ones that I use when I travel overseas: the translation API.

DEMO

Demo:

Part 1: API

Navigate to: <https://cloud.google.com/translate/>

Change English → Spanish and translate “how are you”?

Part 2: G Suite integration

But what about translating a group of sentences? Sure we could invoke the API but let me show you a fun way to invoke this building block right within Google Sheets
<https://docs.google.com/spreadsheets/d/1buZIx1FEIqXRxf3yTCzopoh3z4JebAXhjs54oe5NrM/edit#gid=0>

(would be good to re-create :<https://www.youtube.com/watch?v=MEibCyzPMh4>)

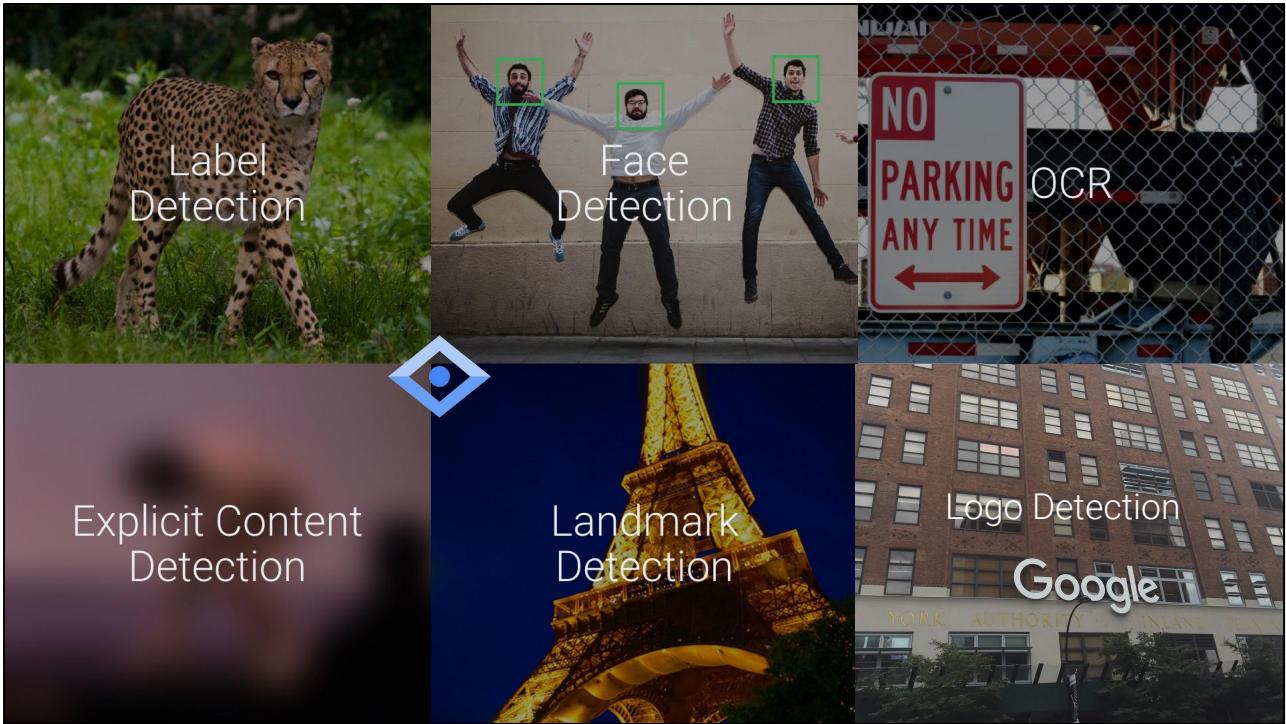
Part 3: Sentiment Analysis

What you saw was just one of the AI building blocks but often we want to do more complex tasks like translate and get the sentiment as well. I'll link cool tech talk from Alicia Williams, a Developer Advocate, where she builds a dashboard on top of the translation and sentiment analysis building blocks. You'll be practicing this in your lab.

<https://support.google.com/docs/answer/3093331?hl=en>

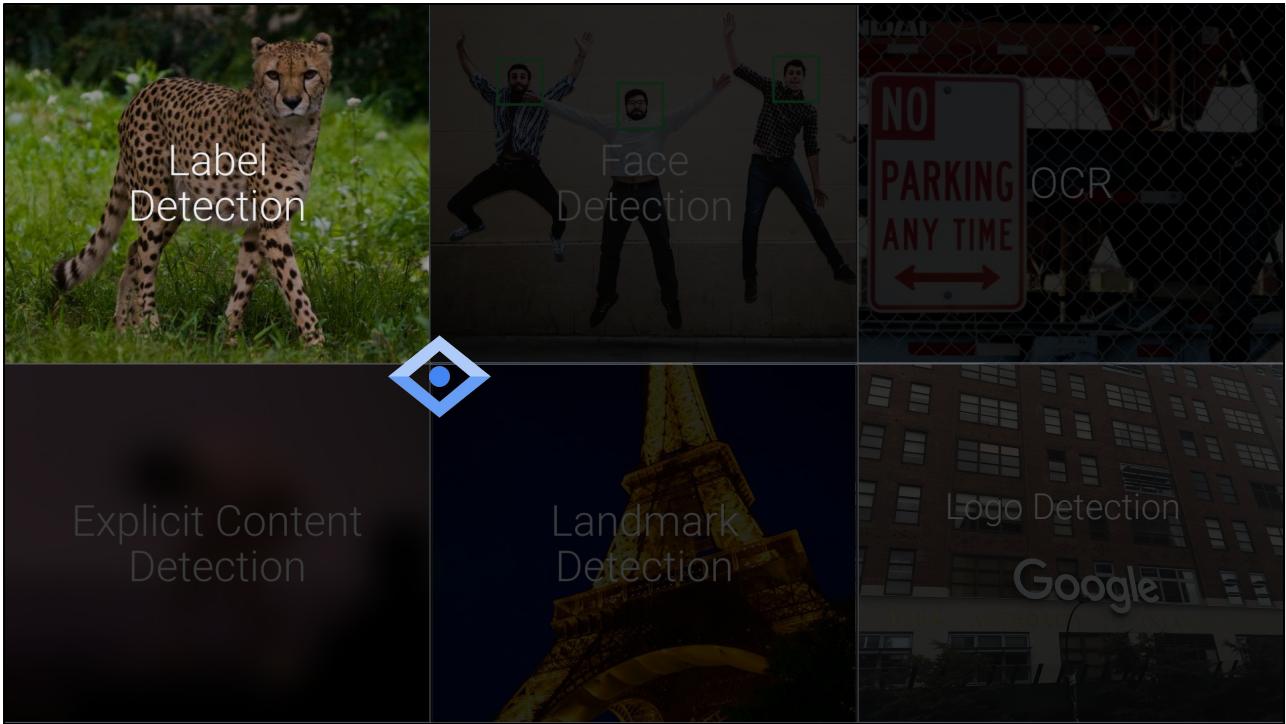
<https://cloud.google.com/blog/products/gcp/analyzing-text-in-a-google-sheet-using-cloud-natural-language-api-and-apps-script>

<https://www.youtube.com/watch?v=Y2wgQjxrPD8&index=73&list=PLBogoxgQVM9v0xG0QTFQ5PTbNrj8uGSS-&t=0s>

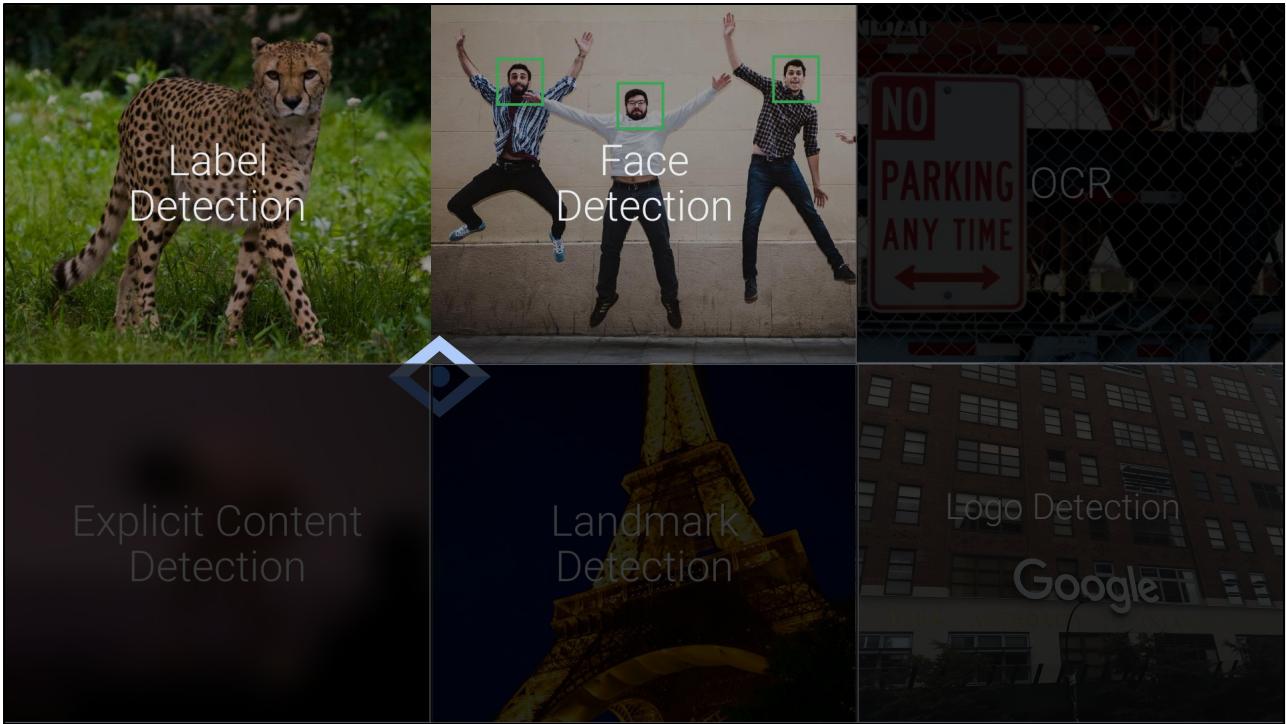


So you saw a demo with text, what about images?

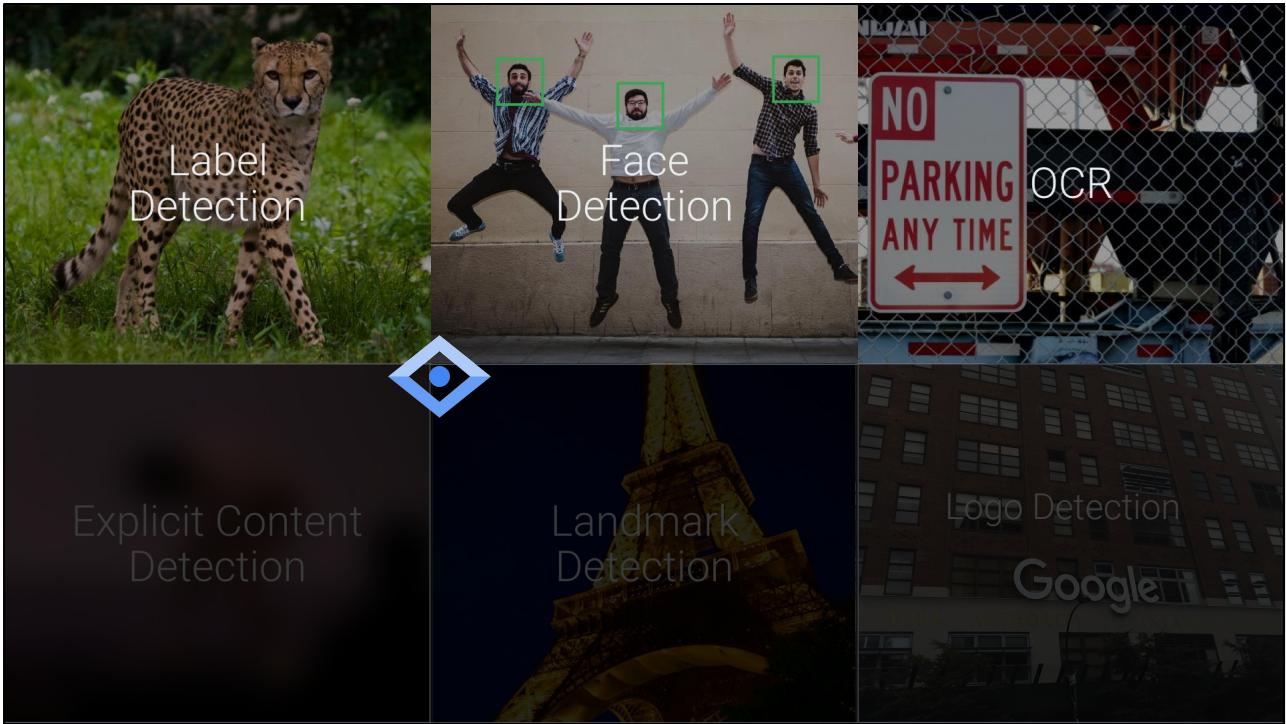
Let me walk you through what Vision API can do:



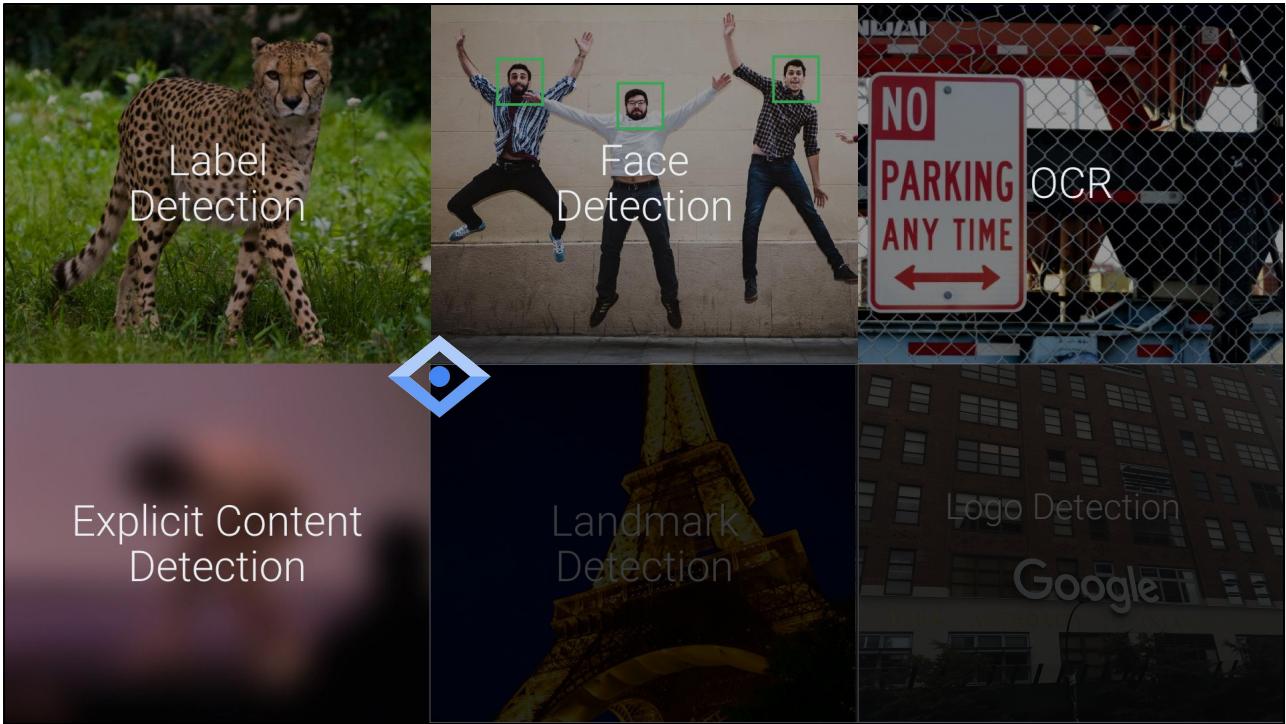
Label Detection: The API can detect broad sets of categories within an image, ranging from modes of transportation to animals.



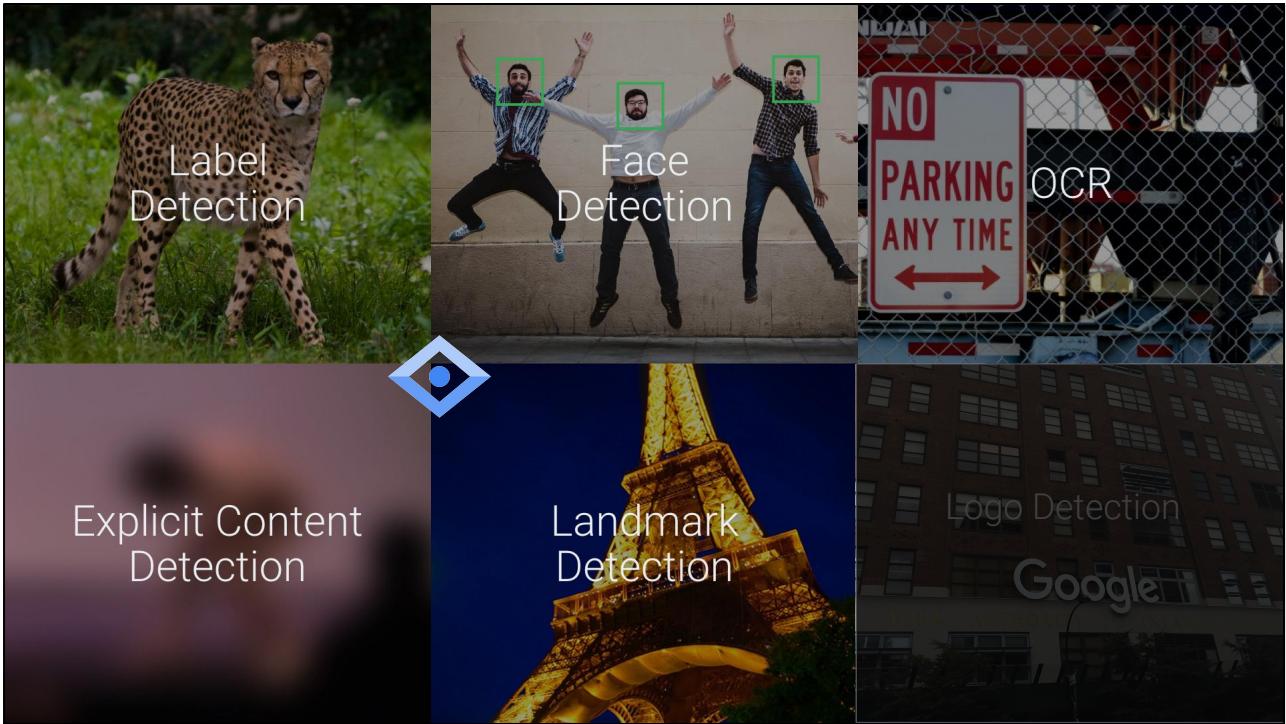
Face Detection: The API can detect multiple faces within an image, along with the associated key facial attributes like emotional state or wearing headwear.



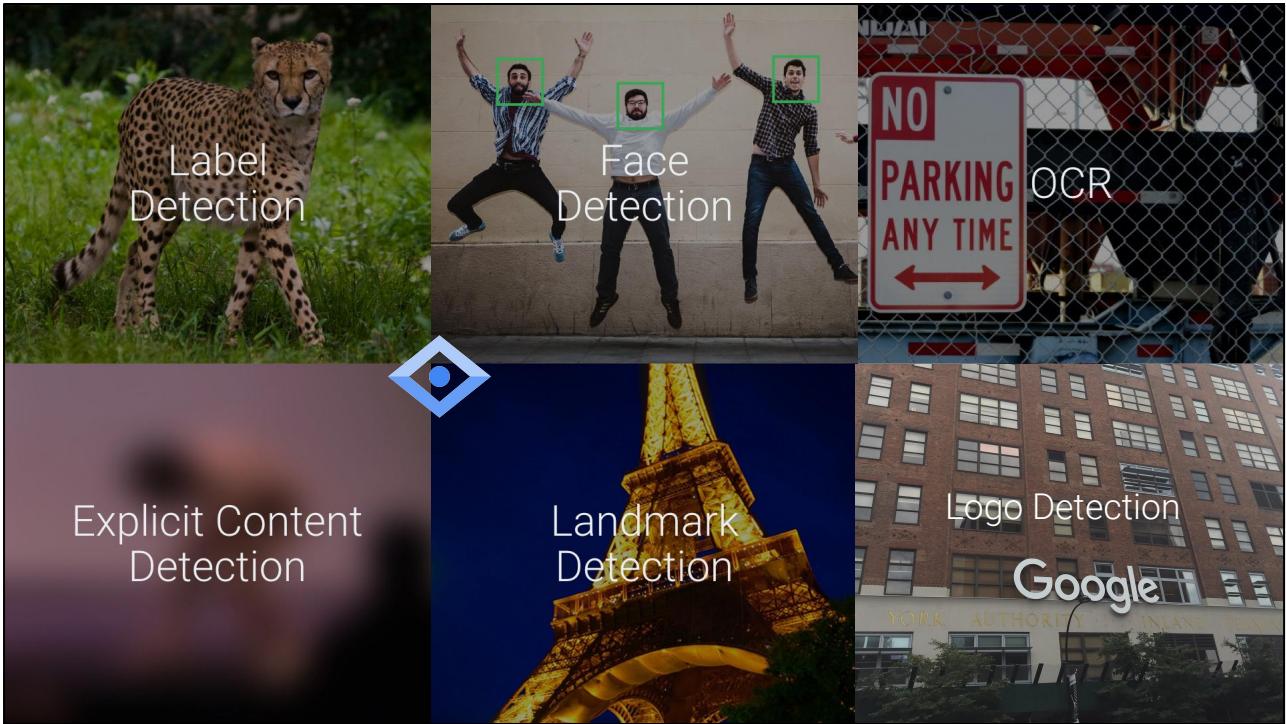
OCR: The API can detect and extract text within an image, with support for a broad range of languages.



Explicit Content Detection: We can detect explicit content like adult content or violent content within an image.



Landmark Detection: The API can detect popular natural and manmade structures within an image.



Logo Detection: We can detect popular product logos within an image.

Face detection

The image shows a composite of two parts. On the left, there is a JSON-like code snippet describing a face annotation for a person taking a selfie. On the right, there is a photograph of a man in a green shirt and a white baseball cap with 'A&M' on it, taking a selfie with a woman in a blue shirt. A blue bounding box surrounds the man's head, and a red bounding box highlights his eyes.

```
"faceAnnotations" : [  
    {  
        "headwearLikelihood" : "VERY_UNLIKELY",  
        "surpriseLikelihood" : "VERY_UNLIKELY",  
        "rollAngle" : -4.6490049,  
        "angerLikelihood" : "VERY_UNLIKELY",  
        "landmarks" : [  
            {  
                "type" : "LEFT_EYE",  
                "position" : {  
                    "x" : 691.97974,  
                    "y" : 373.11096,  
                    "z" : 0.000037421443  
                }  
            },  
            ...  
        ],  
        "boundingPoly" : {  
            "vertices" : [  
                {  
                    "x" : 743,  
                    "y" : 449  
                },  
                ...  
            ]  
        }  
    }  
],  
"detectionConfidence" : 0.93568963,  
"joyLikelihood" : "VERY_LIKELY",  
"panAngle" : 4.150538,  
"sorrowLikelihood" : "VERY_UNLIKELY",  
"tiltAngle" : -19.377356,  
"underExposedLikelihood" : "VERY_UNLIKELY",  
"blurredLikelihood" : "VERY_UNLIKELY"
```



Here's an example of what the JSON response looks like for face detection - it's a picture a Googler on our team took with 2 other teammates on a trip to Jordan.

It returns an object for each face found in an image. You can see as part of the FaceAnnotations some pretty cool attributes like headwearLikelihood and joyLikelihood.

2 **Web annotations**

```
5 {  
6   "entityId": "/m/0gfff2yr",  
7   "score": 5.92256,  
8   "description": "ArtScience Museum"  
9 }  
10 {  
11   "entityId": "/m/016ms7",  
12   "score": 1.44038,  
13   "description": "Ford Anglia"  
14 }  
15  
16  
17
```



```
{  
  "entityId": "/m/0h898pd",  
  "score": 7.4162,  
  "description": "Harry Potter (Literary Series)"  
}
```

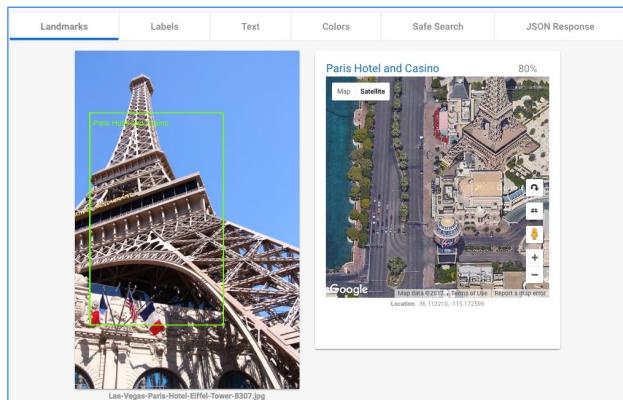
CC-BY 2.0 Rev Stan: <https://www.flickr.com/photos/revstan/6865880240>

The Vision API can also provide annotations on the image by looking up where and in what context this image or visually similar ones have appeared on the web. Take the image of this car which may not be immediately recognizable to some of you. But to others (and to Cloud Vision through web annotations) it's identified as the flying car from the Harry Potter series.

1 Google Cloud

Cloud OnBoard

2 Try it in the browser with your own images



16 cloud.google.com/vision

Take a minute and try Cloud Vision yourself directly in your browser. Navigate to cloud.google.com/vision and scroll down to try the API with your own image to upload.

[cloud.google.com/vision]

The Translation API supports 100+ languages



<https://cloud.google.com/translate/>

Another pre-built ML solution is the Translate API which underlies the product shown above. Simply place your camera over a sign, and it gets auto-translated for you. This is a combination of Vision API (to do optical character recognition) and Translate API (to do actual translation). Vision API supports 90+ languages, can detect the language of source text, and is highly scalable.

You can try this one too on the web.

Wootric uses the Cloud Natural Language API (entity and sentiment) to make sense of customer feedback



Here's an example of a GCP customer who uses the Cloud Natural Language API.

Wootric is a ML-driven customer feedback platform that helps businesses improve their customer service. They collect millions of free text customer survey responses each week. They use the Natural Language API to automate the text processing and sentiment analysis. In this visualization you see the volume of the feedback on the vertical axis and the sentiment on the horizontal axis. Lastly, the coloring of the circles indicates which bucket of feedback that response was automatically classified into (like Usability feedback or Pricing feedback). This allows Wootric and similar organizations to intelligently route and prioritize customer feedback in real time.

<https://cloud.google.com/blog/big-data/2017/03/analyzing-customer-feedback-using-machine-learning>

2
3 When you analyze sentiment, you get a score (positive/negative) as well
4 as a magnitude (how intense?)
5
6

7 The food was excellent, I would definitely go back!

```
8
9 {
10   "documentSentiment": {
11     "score": 0.8,
12     "magnitude": 0.8
13   }
14
15
16
17
18
```

You can try out the pre-built model and get sentiment score for your own free text using the link provided. Note that you'll get a sentiment score (how positive or negative the text was) as well as a magnitude which indicates how intense of a feeling.

The Speech-to-Text API can be used to transcribe audio to text

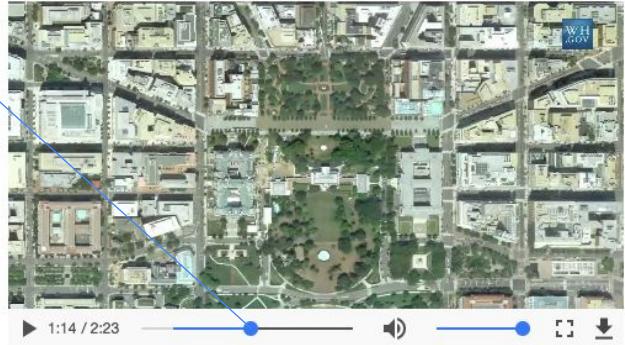


<http://cloud.google.com/speech>

Often multiple pre-built models are used together in a ML system. For example, say if you didn't have free text comments for user reviews but rather had audio from your customer's interactions with your call centers. If you wanted to get the sentiment of the customer's conversation you could first transcribe audio to text with the Cloud Speech API and then use the Natural Language API for sentiment as you saw before.

2
3 Like the Vision API, the Video Intelligence API can identify labels in
4 a video, along with a timestamp

```
5  
6 {  
7     "description": "Bird's-eye view",  
8     "language_code": "en-us",  
9     "locations": {  
10         "segment": {  
11             "start_time_offset": 71905212,  
12             "end_time_offset": 73740392  
13         },  
14         "confidence": 0.96653205  
15     }  
16 }  
17
```



<https://cloud.google.com/video-intelligence/>

One last pre-built model you can leverage is the Video Intelligence API which is similar to the Cloud Vision API except for video instead of images. Here the API can identify labels within a video and when they occurred as well as a confidence level. Here you see at 1:14 in the video the model is 96% confidence that this frame is a “bird’s eye view”

One real customer use case for this API are film companies who are looking to target and recommend movies to an audience based on similar movie-trailer watching behavior. These companies run all their movie trailers through the Video Intelligence API to label key features of the trailer (like “rugged” or “outer space” or “wild west”) and then can programmatically recommend similar movie trailers based on common themes.

Agenda

ML drives business value

How does ML on unstructured data work?

Choosing the right ML approach

- Pre-built AI building blocks
- Using Pre-built AI to create a chatbot
- Customizing Pre-built models with AutoML
- Building a custom model

Demo: Text classification done three ways

Next let's discuss how you can build an intelligent chatbot on top of an already existing solution.

50%

of enterprises will be spending more per annum on bots and chatbot creation than traditional mobile app development by 2021 – *Gartner*

Increasingly, customers do not want to go to your website and click on a button or even send you an email. They want to talk to someone interactively.

The first solution years ago was setting up a call center and manually answering each call but that doesn't scale. So many companies have turned to automation and machine learning to help solve this problem.

Gartner estimates that in a few years, companies will be investing more on creating conversational interfaces like chatbots than even on mobile app development.

Is this chatbot demo just the Speech API paired with a transcription service on top?

Nope, what I'm showing here is a high-level conversational agent tool called Dialogflow which has ML built into its conversations.



Dialogflow

An end-to-end developer platform for building natural and rich conversational experiences

Dialogflow is a platform for building natural and rich conversational experiences. It was formerly named api.ai until Google acquired and rebranded it into Dialogflow.

At its core, Dialogflow is a powerful natural language understanding (NLU) engine to process and understand natural language input. In other words, it lets you easily achieve a conversational user experience by handling the natural language understanding (NLU) for you.



Dialogflow

Dialogflow is an emerging standard for developing conversational interfaces, with a community of **600K+** developers.

Dialogflow is seeing continuous growth in its developer community and is becoming a conversational experience standard.



Dialogflow

Dialogflow users benefit from **Google's** world-class AI assets and capabilities

Google... is an AI company...with a goal to make AI easy, fast, and useful for enterprises and developers.

Dialogflow is built on some of the same world-class AI assets and capabilities that were originally developed for products like Gmail and Search, with new ones being utilized on an ongoing basis. It incorporates Google's ever-growing AI experience, including machine learning expertise, search capabilities, speech recognition, and of course natural language understanding..

Identify key entities in text with Entity Recognition

The screenshot shows a user input field labeled "User says" containing the sentence "weather forecast in San Francisco tomorrow". Below this, a table displays the identified entities and their resolved values:

PARAMETER NAME	ENTITY	RESOLVED VALUE
geo-city	@sys.geo-city	San Francisco
date	@sys.date	tomorrow

Below the table, there are three additional user inputs: "weather for tomorrow", "what is the weather today", and "weather forecast".

It has built-in Entity Recognition enables your agent to identify entities and label by types such as person, organization, location, events, products and media.

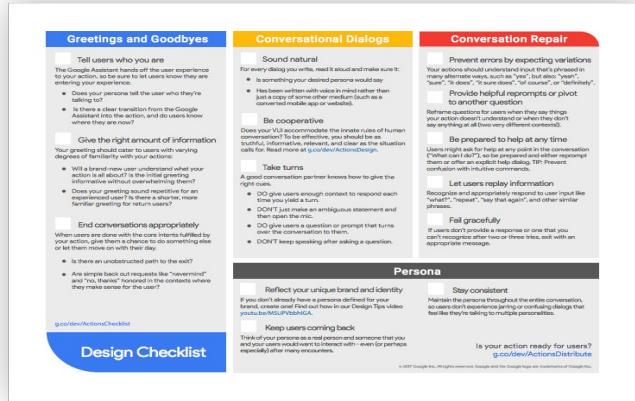
Analyze the overall sentiment
of the conversation



Sentiment
Analysis

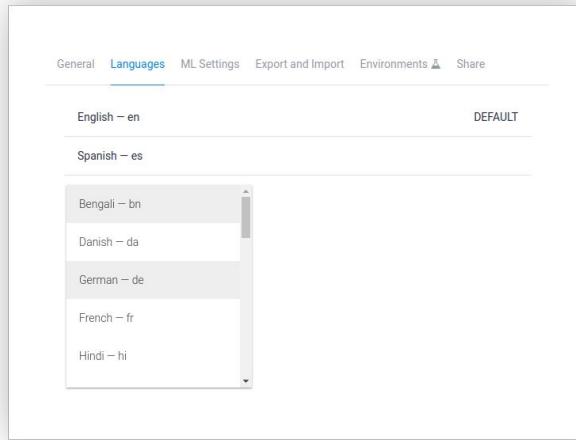
And also Sentiment Analysis to give an understanding of the overall sentiment expressed in a block of text.

Content Classification allows you to classify documents in over 700 predefined categories



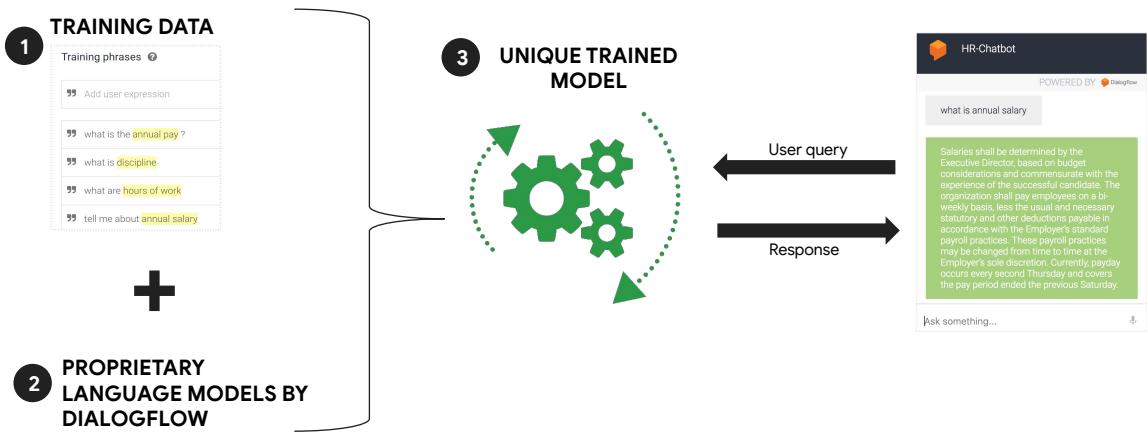
And even Content Classification allows you to classify documents in over 700 predefined categories like common greetings and conversational styles.

Multi-language support



It has Multi-Language support so you can analyze text in multiple languages

How Dialogflow works



Dialogflow works by putting all of these ML capabilities together which you can then optimize for your own training data and use case.

Dialogflow then creates unique algorithms for each specific conversational agent, which continuously learns and is trained and re-trained as more and more users engage with your agent.

Dialogflow benefits for users



Build faster



Engage more efficiently



Maximize reach

The three core benefits to using Dialogflow are:

Building conversational experiences faster, engage end-users more efficiently, and maximize reach across geographies and platforms.

Dialogflow benefits for users



Build faster

- Start training with only a few examples

- 40+ pre-built agents and “small talk” features

Build your agent quickly by starting with just a few training phrases or using one of over 40 prebuilt agents...

Pre-built agents help you get a head start

 Food Delivery Create and manage food and drink orders	 Formats Control default units of measurement	 Hotel Booking Find, create and manage reservations for hotels	 Jokes The agent tells jokes	 Language Settings Set language preferences	 Local Services Search local services and shops
 Navigation Ask for directions	 News Get news stories and manage news feed	 Radio Control playing radio stations	 Reminders Schedule, edit, view and remove reminders	 Maps Search maps for a location	 Music Play and control your music and playlists

These prebuilt agents can be used directly out of the box, or imported into your agent to build on and customize for your own use case. These include everything from food delivery to hotel reservations to news and reminders. You can easily import these prebuilt agents from the Dialogflow console.

Dialogflow benefits for users

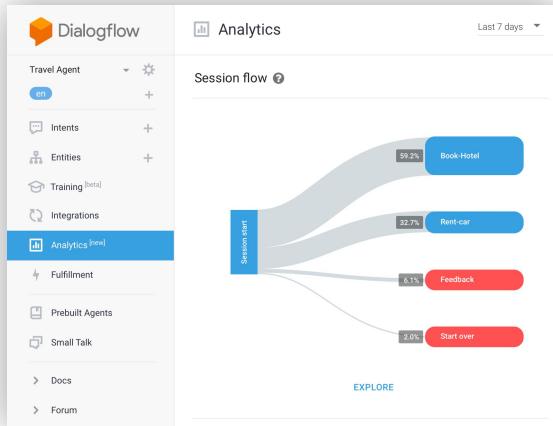


Engage more efficiently

- Built-in, world-class natural language understanding
- Multiple options to connect with backend systems
- Training and analytics

You can rely on Dialogflow's built-in natural language understanding capabilities, model training and analytics that are already taken care of for you.

Analytics offer insights about needed improvements



Dialogflow's built-in Analytics can tell you a lot about users' interactions with your agent. For example, it can show you the breakdown of how often different intents are triggered. This shows you how your users are spending their time in conversation with your agent. This can also help you prune any underutilized intents.

Dialogflow benefits for users



Maximize reach

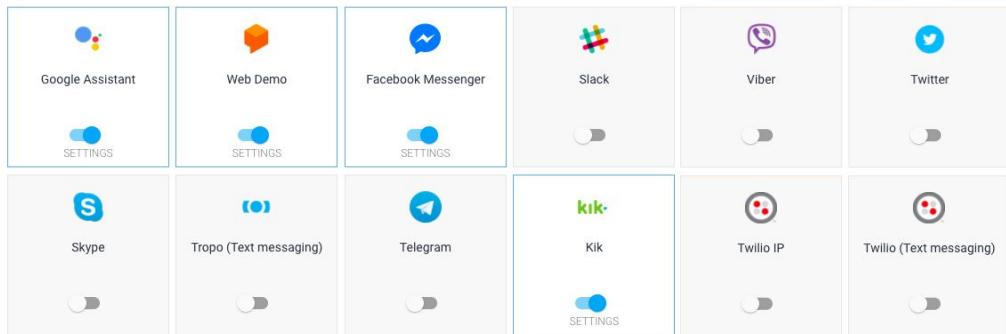
Build once, deploy everywhere

20+ languages supported (most in product category)

14 single-click platform integrations and 7 SDKs

You can also deploy your agent on multiple platforms and utilize Single-click integrations, such as Facebook Messenger, Kik, Twilio, Twitter, Slack, Cisco Spark, etc.

One-click integrations with most major platforms



The Google Assistant integration allows you to quickly deploy your agent to any Assistant-enabled platform, such as the mobile app or a Google Home device. The Web Demo integration allows you to test out your chatbot in an embedded webpage. There are also integrations for other popular apps such as Facebook Messenger, Slack and Twitter, which are enabled with quick authentication.

With all these features and the natural language understanding good built in, lets see the kind of smart, rich, relevant conversations that Dialogflow can bring to life.

Agenda

ML drives business value

How does ML on unstructured data work?

Choosing the right ML approach

- Pre-built AI building blocks
- Using Pre-built AI to create a chatbot
- Customizing Pre-built models with AutoML
- Building a custom model

Demo: Text classification done three ways

We'll continue our tour of pre-built models with high performing custom ML models with one surprising twist -- you can build them with no code. Welcome to AutoML.



Let's say I'm a meteorologist...

Let's say I'm a meteorologist...



I want to predict
weather trends and
flight plans from
images

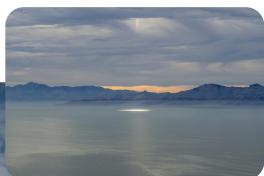
I want to predict weather trends and flight plans from images

There are 10+ different types of clouds



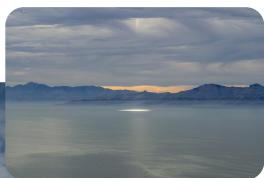
There are more than 10 different types of clouds

There are 10+ different types of clouds



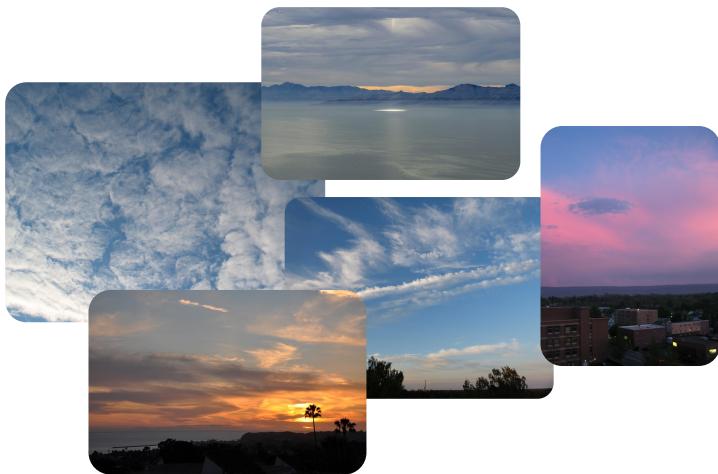
And each type

There are 10+ different types of clouds



Could represent different patterns of weather

There are 10+ different types of clouds



That I may want to know ahead of time before it strikes

There are 10+ different types of clouds



If I'm a pilot or farmer or travel agency or even just someone who wants to know if it's going to rain today or not.

They all indicate different weather patterns



If we want to predict on weather, this means we need to identify not just that there's a cloud, but also what *type* of cloud.

Yes, it's a cloud, but what *type*?



Cumulonimbus

versus



Cirrus

For example, a cirrus cloud is usually associated with fair weather, whereas a Cumulonimbus cloud usually foreshadows rain.

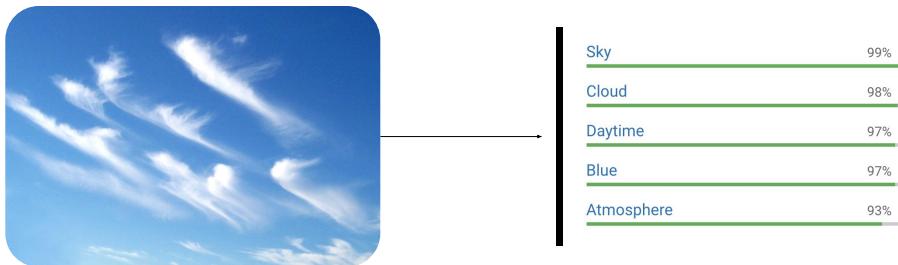
So let's revisit our Vision API and see how well it does on this new problem.

Let's try the Vision API



Here we show the image to the Vision API

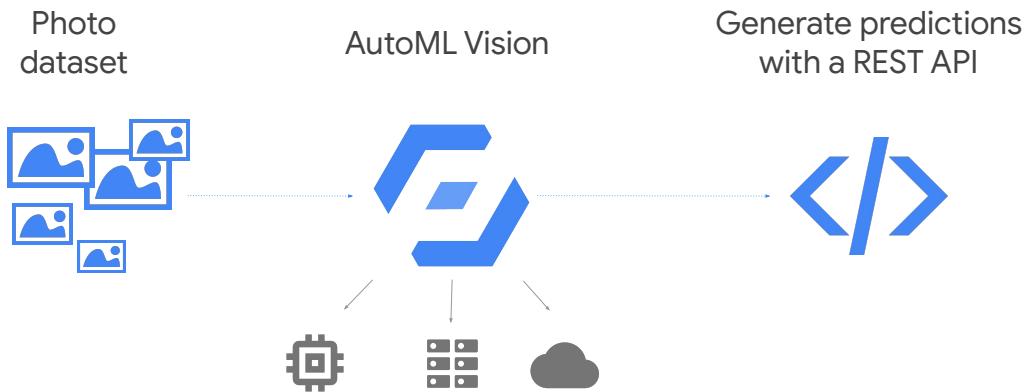
Let's try the Vision API



After I uploaded the image to the Vision API here are the results for the Labels it inferred.

The pre-trained model likely was never taught to recognize cloud types to this granularity. We need something a little more custom that we can train ourselves.

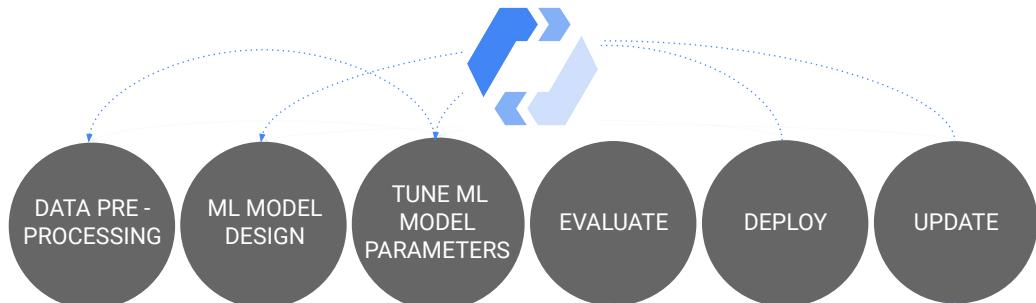
AutoML to the rescue



If you need to extend the capabilities of the AI building blocks, consider AutoML.

Here we will use our own labelled dataset of cloud images and train a custom model using AutoML Vision.

Codeless model building with AutoML



The big takeaway here is not what we did, but what we *didn't* do.

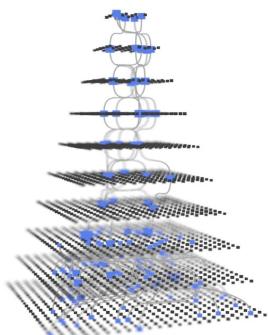
As we talked about at the beginning, typically when building a custom model, you have several steps which are complex and time intensive.

With AutoML, we don't have to do any of that!

However, as great as AutoML is, it can't solve every ML problem which is why you learned how to build a CNN by hand earlier in the course.

Cloud AutoML is built with Neural Architecture Search

Controller: proposes ML models



20K
times

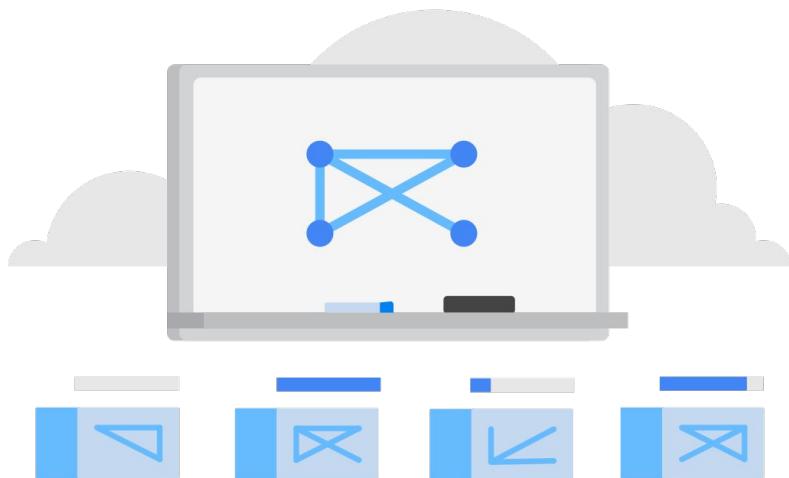
Iterate to
find the
most
accurate
model

Train & evaluate models



Behind-the-scenes, AutoML is powered by the latest ML research. While your model trains, the AutoML platform actually trains and evaluates multiple models and compares them against each other. This NASNet approach or Neural Architecture Search produces an ensemble of ML models and chooses the best one.

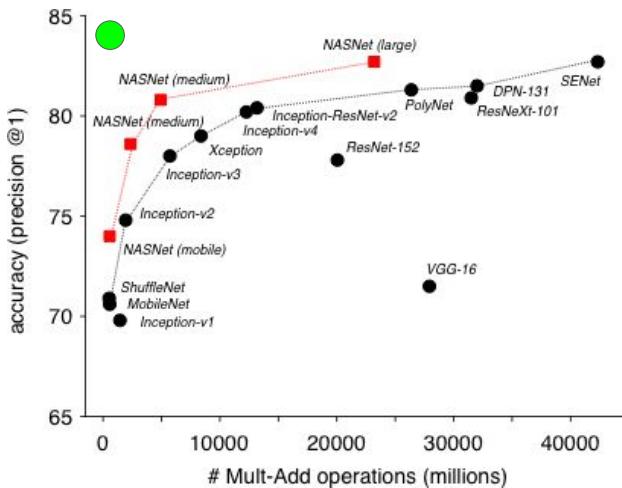
ML models creating ML models



Like Go and self driving cars... deep learning is now doing... deep learning as you saw with the neural architecture search. But how well does a codeless model you build with AutoML compare with some of the other image classification models you may have heard of?

Image: <https://cloud.google.com/automl/>

Inception network for image classification



This graph is a refresh of the best models for ImageNet published in 2017.

The X axis is model size.

The Y axis is accuracy.

The best possible model position would be the green dot. What you want is a small model (left on the X) with great accuracy (top of the Y).

Our world has been trending towards big, heavy, models that aren't exactly like a brain. That's the **black** line.

AutoML, powered by NASNet is shown in the red line. The AutoML networks are smaller and more efficient.

I'll provide a link to the Google AI blog so you can track the latest developments.

<https://ai.googleblog.com/2017/11/automl-for-large-scale-image.html>

AutoML Vision versus Vision API

Attribute	AutoML Vision	Vision API
Objective	Enabling developers with no ML expertise to build state of the art ML models for Images	Enabling ML practitioners to harness power of Google's ML for Images
Primary use case	Classification	Face detection, OCR, Object detection etc.
Data requirements	Images with labelled data	Just Images (may or may not required labelled data)
Output format	Labels with probability	As per the problem
Custom requirements	Can't be customized	Can be used for any custom made solutions
Efforts	Low for solution designing	High for end to end model development
Status	Publicly available	Publicly available

Here's a final recap on when you should use the Vision API vs AutoML Vision.

Recall that with AutoML vision you are primarily doing classification on image data that you are providing to the service. We did this in our demo where we needed the model to learn a label, the type of cloud, that the Vision API didn't know. AutoML requires no coding experience while the Vision API requires you to be familiar with invoking APIs with a programming language of your choice.

Agenda

ML drives business value

How does ML on unstructured data work?

Choosing the right ML approach

- Pre-built AI building blocks
- Using Pre-built AI to create a chatbot
- Customizing Pre-built models with AutoML
- Building a custom model

Demo: Text classification done three ways

Earlier in the course you saw a few ways to create a custom ML model. We'll briefly review them here and point you to additional resources where you can practice building more yourself.

Train and run custom ML in the familiar BigQuery UI



BigQuery ML

One of the easiest ways to create a custom model on a structured dataset from scratch is to try out BigQuery ML. Earlier in the course you practiced:

- Creating a machine learning dataset and identifying features and labels
- Choosing the right model type for your dataset and what you are trying to predict or infer
- Providing any custom model options
- Training the model and evaluating its performance
- Inspecting what the model learned about the weight of each feature
- And predicting on unknown future data

I'll provide additional resources and links for you to practice and learn more about BigQuery ML.

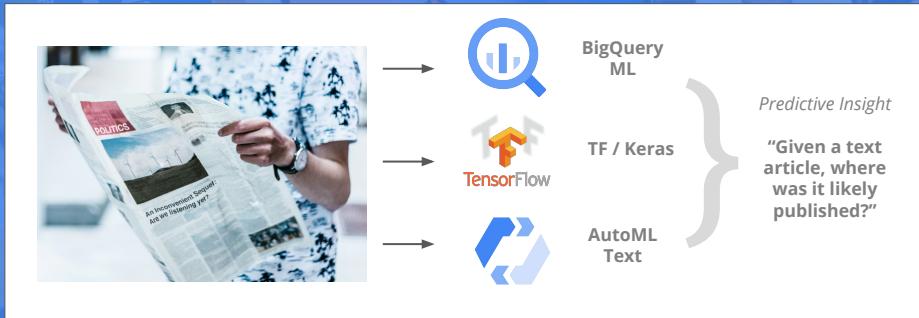
Create custom ML models with TensorFlow



Lastly we mentioned that ML engineers often create their own models using open source libraries like TensorFlow running on GCP. The value of these models can be huge if you build and train them correctly or minimal if they are not done well. If you're looking for experience building TensorFlow models on GCP, checkout our ML on GCP specialization in our additional resources section of the course.

Demo:

Comparing text classification done three ways on Google Cloud Platform



<https://towardsdatascience.com/choosing-between-tensorflow-keras-bigquery-ml-and-automl-natural-language-for-text-classification-6b1c9fc21013>